

HUMAN ARCHITECT · OWL DOSSIER

THE FULFILLMENT LAYER

The managed-AI-employee playbook for UK operators. Model, methodology, objections, case studies, and adoption-ready use cases.

Sell an outcome, not a tool. Charge rent on the loop.

Compiled by OWL | VERIFIED 2026-07-29 | Internal strategy dossier

Live: <https://alkalinear architect.github.io/ai-employee/>

Call: <https://beacons.ai/humanarchitect>

THE MODEL

What a managed AI employee is

A managed AI employee is a deployed AI worker that owns a defined body of work for your business — operated by you, not by the client — for one flat monthly fee. **The client pays to receive an outcome, not to own a tool they have to figure out.**

Price: £5,000 per client, per month. Includes: unlimited agents, unlimited usage, unlimited updates. Managed end-to-end. If it breaks, you fix it fast.

Nick Vasilescu runs this model publicly. Corey Ganim verified it and copied it. Market is early. Margin is ~90%. The moat is not the model — it is the fulfillment layer.

“The agent is leverage. You are still the architect.”

How to close £5k/mo clients

- • **1 · Foot in the door:** lead with a £999 AI audit. Cold prospects say no to £5k. The audit pays you to learn their business AND surfaces the exact pain to solve.
- • **2 · Pick ONE specialised agent** from the audit — the one with the most obvious ROI. One sharp agent beats ten half-built ones.
- • **3 · Price £5,000/mo, all-in.** Unlimited usage, unlimited updates, fully managed.
- • **4 · Ship by day two.** Same-day or next-day delivery changes the entire sales dynamic.
- • **5 · B2B2B retention:** help the client resell the agent to THEIR customers — churn dies.

Speed-to-value rule: the same day a client wires money, ship them something. Agent live by day two.

THE MATH

Why the margin is ~90%

10 clients $\times$ £5,000 = **£50,000/month** recurring.

Overhead per client $\approx$ workspace costs plus modest compute. Self-hosted keeps this tiny.

Key insight: almost no client needs an army of agents — they need one or two very specialised ones. That is why "unlimited" is safe to promise.

"Build once. Fulfill forever. Charge rent on the loop."

What every AI employee needs

- • **1. Computer** — isolated env: VPS/container per client
- • **2. Email** — real inbox the agent can send/receive
- • **3. Phone/IM** — Telegram/Slack (client-facing)
- • **4. Comms channel** — where the client talks to it
- • **5. Tools/connectors** — skills + terminal/file/web
- • **6. Card/payment** — client billing + agent spend rails
- • **7. Knowledge base** — markdown "one fact, one home"
- • **8. Observability** — health checks + Telegram alerts

Miss one and it is a toy, not a worker.

Orchestrator + Swarm

One orchestrator handles the client relationship. Specialised sub-agents execute: outbound, sales follow-up, content, support.

Why split? **Minimise blast radius.** One broken sub-agent must never take down the others.

Rule: never tell the client "the agent builds itself." Say: "I have a build system. You get the agent by Friday."

Repeatable delivery: **Build** → **Onboard** → **Fulfill** → **Pay**. Hands-off every month.

Lead with the KB, not the bot

Before any agent is useful it needs a governed knowledge base. This is the step most operators skip. Don't.

- • Organised by function: Company / Customers / Offers / Sales / Ops / Finance.
- • Every fact carries provenance, owner, sensitivity, review date.
- • Protected tier for credentials; agents get scoped boundaries.
- • Pipeline: raw inputs → wiki → outputs. Maintenance is part of the product.

A governed KB is one of the few moats left in the AI age. Sell it as a service for £5k+.

USE CASES

Adoption-ready client use cases

Real-world use cases mapped to client language and outcome metrics. Each is a standalone conversation starter on a discovery call.

Rule: pick ONE agent per client. One sharp agent beats ten half-built ones.

Exact AI task	Manual pain	Outcome	Monthly value
Pitch deck + brand research + mood board from brief	Designers spend 6–10 hrs/pitch scanning assets and assembling slides	60–75% faster pitch turnaround	£2,800–£4,500
24/7 lead qualification + auto-booking viewings + agent handoff	Agents manually answer DMs; low-quality leads waste 3–5 hrs/day; double-bookings	80%+ leads qualified/booked; no-shows cut 30%	£3,200–£5,000
AI scribe: transcribes consultations, drafts notes, files to records	Clinicians spend 1.5–2 hrs/day typing notes; burnout	Notes in real time; review cut from 90 mins to <15 mins	£3,000–£5,000
Intake agent: pre-qualifies, books discovery calls, intake forms + reminders	Owner handles all scheduling; 30–40% leads go cold before first call	Booking rate lifts 35–55%; admin drops 10–15 hrs/week	£1,800–£3,500
After-hours phone/SMS agent: diagnostics, quotes, books emergency visits	40–60% of callers hit voicemail after hours; lost jobs	60–80% jobs booked without a human; out-of-hours capture 3×	£2,500–£4,500
Document-review agent: reads contracts/invoices, flags anomalies, drafts summaries	Associates spend 40–60% of billable hours on prep; rework from errors	First-pass review cut 70–85%; 10–15 hrs/week to advisory	£4,500–£7,000
Support + CSM agent: tier-1 onboarding, setup walkthroughs, escalation routing	Tickets pile up after launches; CSMs buried; NPS drops	First response <90 sec; 50–65% tickets resolved without human	£2,500–£5,000
Customer-service + returns handler: answers queries, processes returns, upsells	CS team drowned in order-status/return tickets; abandoned carts	60–75% tickets solved without human; AOV lifts 5–12%	£1,800–£3,500
Dispatch + exception agent: auto-assigns drivers, optimises routes, delay alerts	Dispatcher spends all day on rider queries; routing by gut feel; late SLA	Route density +15–25%; failed-delivery rate -20–35%	£3,000–£5,000
Tenant-support agent: answers tenancy questions, logs maintenance, routes contractors	20–40 calls/emails/day per manager; emergency delays; compliance stacks	55–70% tenant queries self-serve; emergency SLA from 4 hrs to 1 hr	£2,800–£4,500
HR agent: policy answers, holiday/OOO tracking, offer/exit letters, compliance	Founder interrupted constantly with policy/admin work; clunky processes	HR admin cut 60–75%; answers instant; compliance risk reduced	£1,500–£2,500

Client-intake + KYC agent: fact-find, Open Banking pull, client folders, gap alerts	3–6 hrs per new client on admin, paperwork, chasing docs	Onboarding from 2 days to same-day; missing-doc turnaround hours not days	£2,500–£4
Reporting agent: pulls Ads/GA4/social APIs, writes narrative reports, anomaly alerts	Account managers spend 6–10 hrs/week on decks; client calls become screenshares	Reporting time cut 80%; calls reallocated to strategy	£2,200–£4
CV screen + candidate-engagement agent: ranks applicants, books interviews, chases feedback	Recruiters spend 60% time reading CVs/admin; candidate ghosting rises	Time-to-shortlist cut 70%; response rate +30–40%	£2,000–£4
Project assistant: status from Jira/Notion, client updates, blocker flags, rebalance	PM spends half the day chasing updates, writing recaps, rescheduling meetings	Status-report time drops 85%; blockers surfaced in hours	£2,000–£3

—

What the research says

UK SMB automation research points to fast payback when the AI owns a defined outcome rather than acting as a chatbot.

- • UK SMBs report saving £500–£2,000/month per focused automation (Insightful AI UK, 57% adoption rate).
- • Professional-services ROI runs 280–370%, with 8–12-month payback (Samyotech, Forrester via Microsoft).
- • Trade/logistics automation typically pays for itself within 2–4 months by preventing one missed job or failed delivery per week.

The common failure mode is shipping a dashboard, not an employee. Outcome ownership is the differentiator.

OBJECTIONS

12 objections + handlers

Objections are requests for reassurance, not rejection. Acknowledge, reframe, close.

Objection	Handler	Close prompt
£5,000 is too expensive	If the task takes 2 hours daily at £40/hr, it costs £2,400/month before mistakes. Owl pays for itself before your first two bad hires.	What does that task cost you right now?
We don't do AI here	Most clients said the same. The agent lives in Slack. You don't touch code. You just watch the work get done.	Can we show you one workflow in 20 minutes?
We already have tools	Tools are not the same as fulfillment. A dashboard is a job. Owl ships the work.	What task is still manual even though you have tools?
We tried AI before	Most AI projects fail because they ship a toy, not a worker. Owl is managed execution, not a demo.	What broke last time — setup, trust, or outcomes?
We're too small	SMBs at £1M–£2M+ pay faster and ask fewer questions. They value outcomes over novelty.	What task costs you 2+ hours every single day?
We need custom	Every agent is custom by default — built from your knowledge, your workflow, your outputs.	Can we scope exactly one role in the next call?
It will replace people	It replaces the manual work nobody wants to do. Your team moves to higher-value work.	What would your team do with 10 hours back this week?
We need to test first	I agree. Start with one agent, one workflow. If it misses the target by day 30, partial refund.	What is the one task you would test first?
Our data is sensitive	Your data stays within your environment. Controlled scope, not blanket access.	Do you have a data policy we should align with?
We need HR approval	Good — HR protects the business. This is an external supplier system, not an internal hire.	Can HR join the scoping call?

It's not a priority right now	That is usually because one painful task is stealing margin and attention.	What task would free you up to focus on growth?
We need to think about it	The only risk is inaction. Margin bleeds while you wait. I will follow up Thursday. Does that work?	What would help you decide by Thursday?

OBJECTIONS

12 objections + handlers

Objections are requests for reassurance, not rejection. Acknowledge, reframe, close.

Objection	Handler	Close prompt
£5,000 is too expensive	If the task takes 2 hours daily at £40/hr, it costs £2,400/month before mistakes. Owl pays for itself before your first two bad hires.	What does that task cost you right now?
We don't do AI here	Most clients said the same. The agent lives in Slack. You don't touch code. You just watch the work get done.	Can we show you one workflow in 20 minutes?
We already have tools	Tools are not the same as fulfillment. A dashboard is a job. Owl ships the work.	What task is still manual even though you have tools?
We tried AI before	Most AI projects fail because they ship a toy, not a worker. Owl is managed execution, not a demo.	What broke last time — setup, trust, or outcomes?
We're too small	SMBs at £1M–£2M+ pay faster and ask fewer questions. They value outcomes over novelty.	What task costs you 2+ hours every single day?
We need custom	Every agent is custom by default — built from your knowledge, your workflow, your outputs.	Can we scope exactly one role in the next call?
It will replace people	It replaces the manual work nobody wants to do. Your team moves to higher-value work.	What would your team do with 10 hours back this week?
We need to test first	I agree. Start with one agent, one workflow. If it misses the target by day 30, partial refund.	What is the one task you would test first?
Our data is sensitive	Your data stays within your environment. Controlled scope, not blanket access.	Do you have a data policy we should align with?
We need HR approval	Good — HR protects the business. This is an external supplier system, not an internal hire.	Can HR join the scoping call?

It's not a priority right now	That is usually because one painful task is stealing margin and attention.	What task would free you up to focus on growth?
We need to think about it	The only risk is inaction. Margin bleeds while you wait. I will follow up Thursday. Does that work?	What would help you decide by Thursday?

Verifiable case studies + quotes

Publicly verifiable operators running managed AI employee / AI employee-for-hire models, with exact quotes.

Source	Claim	Context
Nick Vasilescu / Corey Ganim newsletter	\$5K/mo + fulfillment model	Also confirmed on Build With AI podcast Jul 27 2026; Nick's verified reply: 'thanks for having me on again Corey'
Nick Vasilescu / YouTube (BI-MNjm1tTQ)	Solo operator, managed agents	Demonstrates Dewey build + client deployment walkthrough
Linara Bozieva / Business Insider	27 agents, 5 clients, ~2 hrs/week oversight	Launched after eBay layoff; could scale to 20–25 clients solo; costs under \$1,000/month
Corey Ganim / newsletter (corey-ganim.kit.com)	'10 clients, £50k/month, one person'	Explicit 90% margin math; overhead ~£200 Codex + Orgo workspace per client
Corey Ganim / X post Jun 18 2026	Quantified \$5K retainer justification	'if your time is worth £250/hr and admin is 5 hrs/week, that's £5,000/month you'll buy back'
Boon Media / case studies	Barber: no-shows 18%→4%; SaaS reply 0.8%→6.2%; freight: quote response 45min→5min	Hard metrics from live case-study page
Rebotify / case studies	Energy retailer: drafts days→<1 hr; higher ed: 24hr→1min response; telco: 60–70% automated containment	Enterprise AI employee deployments with public outcomes
Agentive.au	Junior-accountant price benchmark + 24hr deployment	'A junior accountant costs £5,000/month minimum... Deploy in 24 hours... Available 24/7'
Skool / Agent Empire	1.4k members	Free course on building/selling managed agents

No standalone Nick post with the exact '\$5k/mo fulfillment' phrase was found unauthenticated. The quote originates from Corey's recap thread + podcast. Attribute it right.

Your edge: self-hosted

Nick's reference stack is paid. T already owns the free equivalents on this VPS. Same result, higher margin, total control.

Part	Nick pays	T owns free
Computer	Orgo	This VPS, 24/7
Harness	OpenClaw/Hermes	Hermes Agent
Email	agentmail.to	Domain alias
IM	agentphone.ai	Telegram gateway
Tools	Composio	Hermes skills + native tools
Knowledge	fresh vault	~/vault (Obsidian)
Watch	latitude.so	cron → Telegram

Never adopt the paid stack unless T approves a specific upgrade.

Owl execution checklist

- • Foot in the door: offer £999 audit first.
 - • Audit → pick ONE high-ROI agent per client.
 - • Close at £5,000/mo all-in. Agent LIVE BY DAY TWO.
 - • Target SMBs at £1M–£2M+ revenue.
 - • B2B2B retention: help client resell to THEIR customers.
 - • Self-hosted stack; keep overhead ~£200/mo or less.
 - • Template knowledge layer: clients//context.md + discovery-call transcript.
 - • Bonus: sell the AI Second Brain KB service for £5k.
 - • Keep Dewey/Owl income SEPARATE from HumanitAI CIC.
 - • Teach the model publicly on Substack/X to pull leads.
-